

Bovnar — EBNF Grammar

Bovnar (BVNR) v1.1 documentation · EBNF Grammar · 2026-06-18

```

(*) ===== *)
(*) BOVNAR - EBNF derived from parser/lexer/validator implementation *)
(*) Version 1.1 Notation: ISO/IEC 14977:1996 *)
(*) { } = zero or more [ ] = optional *)
(*) ( ) = grouping | = alternation *)
(*) , = concatenation ; = rule terminator *)
(*) "..." = literal terminal (*...) = comment *)
(*) Updated: 2026-06-17 - drop the residual "1.1 (draft)" label (spec *)
(*) 1.1 is released); note datetime also accepts a width parameter; *)
(*) reframe base-unit as a lookup into the full 163-unit / 487-spelling *)
(*) table (NOT the closed SI subset it previously claimed). *)
(*) 2026-06-16 - consistency pass vs the implementation: *)
(*) allow ws/comments before the assignment ";" and the "/" array *)
(*) separator; move the inline unit to a value-only scalar-with-unit *)
(*) rule so array elements are correctly excluded; bound \u{...} to 1-6 *)
(*) hex digits and version-int to no-leading-zero; inline type-spec; *)
(*) split the array-homogeneity / duplicate-key constraints into the *)
(*) LEXER/VALIDATOR tier (the streaming reader) vs the DOM-BUILDER *)
(*) tier (only a DOM parse rejects ragged siblings, mixed kinds, and *)
(*) duplicate keys; the streaming reader accepts them). *)
(*) 2026-06-15 - spec 1.1 (additive over the frozen 1.0): *)
(*) version directive (!bovnar M.N), datetime type family, *)
(*) \xHH / \u{...} string escapes, reference array indexing. *)
(*) Earlier (2026-05-29): null/true/false/on/off keywords, bool *)
(*) type family, special numbers respelled nan / inf / ninf *)
(*) ===== *)

(*) ----- *)
(*) 1. Top-level stream *)
(*) ----- *)

(*) Valid end-of-stream states (lexer states where EOF is accepted): *)
(*) undefined - nothing has been read yet (empty stream) *)
(*) first_comment_intro - inside a leading comment that runs to EOF *)
(*) with no trailing newline (comment-only stream) *)
(*) first_bom - BOM or first-line comment (its newline seen), *)
(*) no assignment yet *)
(*) value_outro - the semicolon of the last assignment consumed *)
(*) In all four cases struct_nesting_level must be zero. *)

(*) Events emitted to the validator callbacks (bvnr_event_t): *)
(*) *)
(*) ev_stream_start - once at the very start of bvn_lex_run *)
(*) ev_stream_end - once at the very end of bvn_lex_run, *)
(*) after the last value's terminator is consumed *)
(*) ev_assignment_start - on the identifier token of an assignment *)
(*) ev_struct_start - on "{" *)
(*) ev_struct_end - on "}" *)
(*) ev_array_row_start - on "[" opening a bracket row *)
(*) ev_array_row_end - on "]" closing a bracket row *)
(*) ev_array_dim_start - on "/" between rows; fired immediately *)
(*) when "/" is consumed, before the next *)
(*) "[" opens the next row *)
(*) ev_octet_stream_start - on the NUL byte that enters binary mode *)
(*) ev_octet_stream_end - on the 0x00 tag ending an octet stream *)
(*) *)
(*) For every typed value (including values whose type annotation is *)
(*) synthesised by the validator - see §4) the following sequence *)
(*) is emitted before ev_data: *)
(*) *)
(*) ev_type_annotation_start - once, before params *)
(*) ev_type_annotation_type_family - the family name *)
(*) ev_type_annotation_type_family_parameter - zero to three *)
(*) times, once per present parameter (width, *)
(*) base, unit), in that order *)

```

```

(*      ev_type_annotation_end          – once, after params      *)
(*      ev_data                        – the value token          *)
(*)
(*)
(*) ev_type_annotation_start is emitted to on_unverified first    *)
(*) (raw bytes only), then to on_verified (with value_type and    *)
(*) value_unit filled in). All other type-annotation events and   *)
(*) ev_data are emitted to both callbacks simultaneously.        *)
(*) For bool keywords (true/false/on/off) the same default type- *)
(*) annotation sequence is emitted, synthesising <bool> when no   *)
(*) explicit annotation was given (no width/base/unit params).   *)
(*)                                                                *)
(*) ev_data is also emitted for null values, symbols, references, *)
(*) and octet-stream chunks without a preceding type-annotation  *)
(*) sequence. Struct values do NOT produce ev_data; they are     *)
(*) signaled by ev_struct_start and ev_struct_end.               *)

(*) Only blank characters – NOT comments – may precede the version *)
(*) directive: it is recognised solely as the very first comment, *)
(*) earlier comment would demote it to an ordinary one (see below). *)
stream      = [utf8-bom] , {ws-char} , [version-directive] , ws ,
              {assignment , ws} ;

(*) UTF-8 BOM: legal only as the very first three bytes of the stream. *)
utf8-bom    = "\xEF\xBB\xBF" ;

(*) Version directive (spec 1.1): an OPTIONAL leading comment of the *)
(*) form "#!bovnar <major>.<minor>" declaring the spec version the *)
(*) document targets. It is recognised only as the very first comment *)
(*) (after the optional BOM and whitespace); a "#!bovnar ..." on any *)
(*) later line is an ordinary comment. Lexically it IS a comment, so a *)
(*) spec-1.0 reader skips it transparently – the declaration is purely *)
(*) additive. A 1.1+ reader parses it, exposes the version, and (in *)
(*) strict mode) rejects a version it does not support. Each component *)
(*) is a decimal integer with no leading zero (a bare "0" excepted). *)
(*) A malformed version after the "!"bovnar" prefix is *)
(*) error_invalid_spec_version. *)
(*) The directive accepts ONLY horizontal whitespace (HT or SP) around *)
(*) its components – not the LF/CR/VT/FF that "ws-char" also covers. *)
version-directive = "#" , "!"bovnar" , hspace , {hspace} ,
                  version-int , "." , version-int ,
                  {hspace} , ("x0D" | "x0A" | (* eof *)) ;
hspace           = "\x09" (* HT *) | "\x20" (* SP *) ;
(*) version-int: a decimal integer with no leading zero (a bare "0" is the *)
(*) only form starting with "0"); each component must fit in a uint16. *)
version-int      = "0"
                  | ( ("1"|"2"|"3"|"4"|"5"|"6"|"7"|"8"|"9") , {DIGIT} ) ;

(*) Whitespace: blank characters and comments, freely intermixed. *)
ws              = {ws-char | comment} ;

(*) At least one whitespace character, followed by optional further *)
(*) whitespace and comments. Used as the mandatory separator before *)
(*) an inline unit suffix. *)
ws-mandatory    = ws-char , {ws-char | comment} ;
ws-char         = "\x09" (* HT *)
                  | "\x0A" (* LF *)
                  | "\x0B" (* VT *)
                  | "\x0C" (* FF *)
                  | "\x0D" (* CR *)
                  | "\x20" ; (* SP *)

(*) Comment: "#" through the next LF, CR, or end of input. *)
(*) Control bytes 0x00–0x08, 0x0E–0x1F, and 0x7F are hard errors *)
(*) inside a comment body. *)
comment         = "#" , {comment-char} , ("x0D" | "x0A" | (* eof *)) ;
comment-char    = "\x09" (* HT *)
                  | "\x0B" | "\x0C" (* VT, FF *)

```

```

    | printable-byte          (* 0x20-0x7E *)
    | high-byte ;            (* 0x80-0xFF *)

(* ----- *)
(* 2. Assignments *)
(* ----- *)

(* The leading "." is a syntactic sigil; it is not stored in the key. *)
(* At least one id-start character after "." is required; *)
(* "." is error_empty_identifier. *)
(* Comments are permitted between the last key character and "=". *)

assignment = "." , key , ws , "=" , ws , value , ws , ";" ;

key = id-start , {id-body-char} ;

(* First character of an identifier, symbol, or reference segment: *)
(* A-Z, a-z, "_", and UTF-8 leader bytes 0xC3-0xF4. *)
(* 0xC2 is explicitly rejected (ACT_NONE) at all token-start *)
(* positions, so code points U+0080-U+00BF cannot begin these tokens. *)
id-start = ALPHA | "_" | utf8-leader ;

(* Body characters of an identifier: *)
(* Accepted: "+", "-", digits, alpha, "_", *)
(* UTF-8 leader bytes 0xC3-0xF4, UTF-8 continuation bytes 0x80-0xBF. *)
(* Hard-error bytes: 0x00-0x1F, 0x7F, 0xC2, and ASCII punctuation: *)
(* " # , . / ; < > [ ] { } = ! $ % & ' ( ) * : ? @ \ ^ ` | ~ *)
(* "=" terminates the body (ACT_value_intro). *)
(* Whitespace terminates the body (ACT_to_identifier_outro), after *)
(* which only "=" is accepted (with comments permitted between). *)
id-body-char = "+" | "-"
    | DIGIT
    | ALPHA
    | "_"
    | utf8-leader (* 0xC3-0xF4 *)
    | utf8-continuation ; (* 0x80-0xBF *)

(* ----- *)
(* 3. Values *)
(* ----- *)

(* A value is the right-hand side of an assignment or an element *)
(* inside an array row; the optional type annotation applies in both *)
(* contexts. *)
(* *)
(* Typed null elements: in array context, a type annotation not *)
(* followed by any raw value token (i.e. the next byte is "," "]" *)
(* or ";") is a null value carrying the annotated type. This is *)
(* handled by ACT_type_null_then_new_array_value, *)
(* ACT_type_null_then_array_outro, and ACT_type_null_then_value_outro *)
(* from the type_outro state. *)

value = [type-annotation , ws] , ( scalar-with-unit | raw-value ) ;

(* A bare number or string scalar may carry an inline unit suffix – but ONLY *)
(* at value position, never inside an array element (array-elem uses raw- *)
(* value, which has no unit option). The lexer reaches the inline-unit state *)
(* only from number_outro / string_outro, so no other value form takes one. *)
scalar-with-unit = ( number | string ) , ws-mandatory , inline-unit ;

raw-value = null-value
    | bool-value
    | special-number
    | datetime-literal (* spec 1.1 *)
    | number

```

```

| string
| symbol
| reference
| array
| struct
| octet-stream ;

(* datetime-literal (spec 1.1): an ISO-8601 UTC timestamp accepted in *)
(* place of the integer epoch-seconds carrier. The lexer accumulates the *)
(* shape (dtlit_* states, entered from a digit run followed by "-"); the *)
(* validator strictly range-checks the fields and converts the UTC civil *)
(* instant to the integer carrier on the declared epoch (unix when a bare *)
(* literal infers <datetime:64,unix>). Whole-second carrier: a frac- *)
(* tional part (any digit count) does not change the value but is kept *)
(* verbatim and re-emitted so the literal round-trips. A "+HH:MM" / *)
(* "-HH:MM" offset is folded to true UTC before conversion. *)
(* The atomic GNSS epochs (gps/galileo/glonass/beidou) reject a literal. *)
datetime-literal = year , "-" , month , "-" , day ,
                  [ "T" , hour , ":" , minute , ":" , second ,
                    [ "." , dt-fraction ] , [ dt-zone ] ] ;
dt-fraction      = digit , { digit } ; (* kept verbatim; carrier is whole seconds *)
dt-zone          = "Z" | ( "+" | "-" ) , hour , ":" , minute ; (* tz offset *)
year             = digit , digit , digit , digit ;
month            = digit , digit ; (* 01..12 *)
day              = digit , digit ; (* 01..28/29/30/31 per month *)
hour             = digit , digit ; (* 00..23 *)
minute           = digit , digit ; (* 00..59 *)
second           = digit , digit ; (* 00..59 *)

(* null-value: either the reserved keyword "null" or the empty *)
(* production. *)
(* At assignment level: ".foo = ;" and ".foo = null;" are equivalent. *)
(* In array context: a leading or trailing comma, two consecutive *)
(* commas, a bare "null", or a type annotation immediately followed *)
(* by "," / "]" all denote a null element. *)
null-value       = "null" | (* empty *) ;

(* bool-value: a reserved keyword in value position. "on" is an *)
(* alias for "true" and "off" for "false"; the validator collapses *)
(* all four spellings to a vt_bool value (canonical text "true" / *)
(* "false") and synthesises a <bool> type annotation when none is *)
(* given (see §4). An explicit <bool> annotation accepts only these *)
(* keywords as its value. *)
(* *)
(* Lexically these keywords – together with "null" – are ordinary *)
(* symbol tokens (§7); the validator reserves the exact words and *)
(* reclassifies them. A longer word that merely starts with one *)
(* (e.g. "truthy", "nullable") remains a plain symbol. *)
bool-value       = "true" | "false" | "on" | "off" ;

(* _____ *)
(* 4. Type annotations *)
(* _____ *)

(* Lexer: a keyword state machine (tf_* states) recognises the eight *)
(* family names via six keyword paths (float_fix / float_dec share *)
(* the "float" prefix; see §4 below). After the family keyword the *)
(* raw bytes of the ":"... *)
(* parameter string are accumulated into type_data by copy_type_byte. *)
(* *)
(* Validator: bv_n_parse_type_annotation parses the accumulated string *)
(* and the validator emits the structured type-annotation event *)
(* sequence described in §1. *)
(* *)
(* Default synthesis: when a number or string value carries no *)
(* explicit type annotation (value_type is still vt_plain), the *)

```

```

(* validator synthesises a default before emitting ev_data: *)
(* strings / special-numbers → <utf8> / <float:64,_10,no_unit> *)
(* floats (has "." or "e"/"E") → <float:64,_10,no_unit> *)
(* negative integers → <sint:64,_10,no_unit> *)
(* plain integers → <uint:64,_10,no_unit> *)
(* bool keywords (true/false/on/off) → <bool> *)
(* The full type-annotation event sequence is always emitted for *)
(* number and string values, whether annotated explicitly or not. *)

type-annotation = "<" , ws , param-type , ws , ">" ;

(* Types with an optional parameter list after ":". *)
(* Parameters are identified by class and may appear in any order; *)
(* at most one of each class (width, base, q, unit) is permitted. *)
(* *)
(* "utf8" goes to type_body_outro after the keyword; the lexer *)
(* accepts ":" and the parameter bytes, but utf8 is a parameterless *)
(* family – any width, base, q, or unit parameter is rejected by *)
(* bv_n_parse_type_annotation as error_illegal_value_type (see bool). *)
(* *)
(* "float_fix" and "float_dec" are not separate keyword paths in the *)
(* lexer state table. The lexer fires ACT_tf_float_done after the *)
(* shared "float" prefix (states tf_f → tf_fl → tf_flo → tf_floa *)
(* → ACT_tf_float_done), which stores "float" in type_data and *)
(* transitions to type_body_outro. The subsequent bytes "_fix" or *)
(* "_dec" are accumulated via copy_type_byte. The final string *)
(* ("float", "float_fix", or "float_dec") is dispatched in *)
(* bv_n_parse_type_annotation. *)
(* "bool" and "utf8" are parameterless families: any width, base, q, *)
(* or unit parameter is error_illegal_value_type. bool's only valid *)
(* values are the bool keywords (see bool-value, §3). *)
(* "datetime" (spec 1.1) is a timestamp: a signed integer seconds count *)
(* validated like sint. It accepts an optional width parameter (as sint) *)
(* plus one epoch-name parameter (unix|tai|gps|mjd|ntp|galileo|glonass| *)
(* y2000|beidou; default unix), e.g. <datetime:64,unix>. A numeric base, *)
(* q, or unit parameter is error_illegal_value_type. Available only when *)
(* the document declares "#!bovnar 1.1" or newer. *)
param-type = ("uint" | "sint" | "float" | "float_fix" | "float_dec" | "utf8"
              | "bool" | "datetime")
            , [ws , ":" , ws , type-param-list] ;

(* The parameter list is accumulated byte-by-byte into type_data by *)
(* the copy_type_byte / type_body_outro states, and later re-parsed *)
(* by bv_n_parse_type_annotation. *)
(* *)
(* Bytes accepted in the accumulated string: *)
(* A-Z, a-z, 0-9, "*", "+", ",", ".", "/", "-", ":", "^", "_" *)
(* UTF-8 continuation bytes 0x80-0xBF *)
(* UTF-8 leader bytes 0xC2-0xF4 (0xC2 IS accepted here, unlike *)
(* identifier / symbol / reference token starts) *)
(* Whitespace is accepted but not accumulated; it transitions to *)
(* type_body_outro, which accepts the same bytes as copy_type_byte *)
(* plus "#" (0x23, entering comment_intro), while treating whitespace *)
(* as ignored rather than re-triggering the outro transition. *)
(* ">" always terminates the annotation. *)
type-param-list = type-param , {ws , "," , ws , type-param} ;

type-param = width-param (* e.g. 32, 64 *)
            | base-param (* e.g. _16, _64, _85 *)
            | q-param (* e.g. q8, q16 – float_fix only *)
            | unit-param ; (* e.g. m~s, Gi~B, m^2 *)

(* Width: one or more decimal digits. 0 means "no width restriction".*)
(* float: 0, 16, or any multiple of 32 up to BVN_FLOAT_MAX_PREC.*)
(* float_fix, *)
(* float_dec: 0, 16, 32, 64, 128, 256. *)
width-param = DIGIT , {DIGIT} ;

```

```

(* Base: underscore followed by one or more decimal digits. *)
(* Valid values: 2-62, 64 (standard Base64), 85 (Ascii85). *)
(* Bases 64 and 85 are uint-only: their alphabets use '+'/'-' as *)
(* digits, leaving no sign character, so sint+64/85 is illegal. *)
(* "float" accepts only base 10 (or absent, → 10) or base 16; *)
(* other bases → error_illegal_value_type. *)
(* "float_fix" and "float_dec" reject the base parameter entirely: *)
(* any base-param → error_illegal_value_type. *)
(* A non-decimal base with uint/sint should use a quoted string value; *)
(* a bare literal is not rejected by the validator but will be parsed *)
(* as a symbol token, causing a type/value mismatch instead. *)
base-param = "_" , DIGIT , {DIGIT} ;

(* Q-parameter: lowercase "q" followed by one or more decimal digits. *)
(* Specifies the number of fractional bits for the float_fix Q-format. *)
(* Value = raw_signed_integer × 2^(-Q). *)
(* Valid only for "float_fix"; rejected for all other families. *)
(* Valid range: 0 ≤ Q < effective_width (width 0 → effective 64). *)
(* Q is stored in value_type_spec_t.base and retrieved via *)
(* bvn_effective_q; bvn_effective_base always returns 10 for *)
(* float_fix and float_dec. *)
q-param = "q" , DIGIT , {DIGIT} ;

(* Unit parameter: everything up to the next "," or ">". *)
(* Parsed against the unit table by bvn_parse_unit. *)
(* "no_unit" is a reserved keyword meaning explicitly dimensionless. *)
(* Valid for: uint, sint, float, float_fix, float_dec. *)
unit-param = unit-char , {unit-char} ;
unit-char = (* any byte accepted by copy_type_byte that is not *)
            (* ",", (0x2C) or ">" (0x3E): *)
            (* A-Z, a-z, 0-9, "*", ".", "/", "-", ":", "^", "_", *)
            (* "~", "$", "(", ")", or a UTF-8 multi-byte byte *)
            (* (0x80-0xF4) *) ;

(* — Unit sub-grammar (semantic, enforced by bvn_parse_unit) — *)
(* *)
(* resolved-unit = "no_unit" | unit-expr ; *)
(* unit-expr = unit-factor , {unit-sep , unit-factor} ; *)
(* unit-factor = unit-component | "(" , unit-expr , ")" ; *)
(* *)
(* unit-sep = "*" (* product – U+002A *) *)
(* | "." (* product – U+00B7 *) *)
(* | "/" ; (* division *) *)
(* "." is the middle dot, encoded as 0xC2 0xB7 (UTF-8 for U+00B7). *)
(* "*" and "." are semantically equivalent (multiplication). *)
(* *)
(* Semantics of "/": the first "/" switches every subsequent factor *)
(* into the denominator (its stored exponents are negated). A "(...)" *)
(* group is a sub-expression evaluated independently; when it falls in *)
(* the denominator its net exponents are negated as a whole, so *)
(* "k~g/(m·s²)" = kg·m-1·s-2 and "(k~g/m)·s²" = kg·m-1·s². Additional *)
(* "/" separators at one level never toggle back to the numerator. *)
(* An explicit separator is required before a group ("m·(s)", not *)
(* "m(s)"); a group is not (yet) followed by its own exponent; paren *)
(* nesting is bounded at 16. *)
(* *)
(* Empty components or groups (e.g. "m//s", "()", "m*s"), unmatched *)
(* parentheses, and more than BVNR_MAX_UNIT_COMPONENTS = 8 total *)
(* components per unit string all produce error_unit_illegal. *)
(* *)
(* unit-component = [prefix "~"] base-unit [unit-exponent] ; *)
(* *)
(* When a prefix is present the separator "~" is mandatory. *)
(* A bare base-unit with no prefix requires no separator. *)
(* *)
(* si-prefix = "Q"|"R"|"Y"|"Z"|"E"|"P"|"T"|"G"|"M"|"k"|"h"|"da" *)

```



```

(*      | "d"|"c"|"m"|"μ"|"n"|"p"|"f"|"a"|"z"|"y"|"r"|"q" ; *)
(*      (μ = U+00B5: 0xC2 0xB5) *)
(*)
(*)
(*) iec-prefix = "Ki"|"Mi"|"Gi"|"Ti"|"Pi"|"Ei"|"Zi"|"Yi"|"Ri"|"Qi" ; *)
(*) Ri (2^90) and Qi (2^100) were added in IEC 80000-13:2022 ed. 2. *)
(*)
(*) base-unit = a symbol looked up in the fixed unit table by *)
(*) bvn_parse_unit (src/utils/bovnar_si_units.c): 163 named units with *)
(*) 487 accepted spellings (canonical symbols plus long-form and plural *)
(*) aliases), spanning SI, imperial / US customary, CGS, radiation, *)
(*) surveying, culinary, Old German and digital-storage quantities. *)
(*) Doc 2 (the unit-system spec) is the authoritative registry; the SI *)
(*) coherent / common subset below is REPRESENTATIVE only, NOT a closed *)
(*) alternation: *)
(*) b B s m g A K mol cd Hz N Pa J W V Ω F C S Wb T H lm lx Bq Gy Sv *)
(*) kat rad sr L l min h d wk yr °C ° degC degrC degrees degree degr *)
(*) deg t bar eV Da au ha *)
(*) Symbols are case-sensitive and a few spellings collide across roles, *)
(*) resolved by context: "P" = poise (base unit) vs peta (SI prefix), *)
(*) "G" = gauss vs giga. (Ω = U+2126: 0xE2 0x84 0xA6; ° = U+00B0: 0xC2 *)
(*) 0xB0) *)
(*)
(*) unit-exponent = [exp-sign] exp-digit *)
(*)      | "^" ["-"|"+"] ASCII-digit ; *)
(*) *)
(*) exp-sign = "+" (* U+207A: 0xE2 0x81 0xBA, positive / no-op *) *)
(*)      | "-" ; (* U+207B: 0xE2 0x81 0xBB, negate exponent *) *)
(*) *)
(*) exp-digit = "1"|"2"|"3"|"4"|"5"|"6"|"7"|"8"|"9" ; *)
(*)      (U+00B9 / 00B2 / 00B3 / 2074-2079) *)
(*) *)
(*) "+" is accepted but has no effect (same as absent). *)
(*) The ASCII form "[+-]?[1-9]" maps to the same internal constants *)
(*) as the Unicode superscript form. "^0" is not a valid exponent. *)

(*) ----- *)
(*) 5. Numbers *)
(*) ----- *)

(*) The lexer is digit-agnostic: any ASCII digit 0–9 is accepted in a *)
(*) bare number literal without base validation. Digit validity against *)
(*) a non-decimal base is checked later by bvn_acc_parse_number. *)
(*) Leading zeros are not rejected by the lexer ("007" is valid). *)
(*) Only "e"/"E" are accepted as exponent markers in bare literals. *)
(*) "p"/"P" (binary exponent) is recognised by the validator inside a *)
(*) quoted-string number value of a base-16 float, e.g. *)
(*) <float:64, 16> "1.8p+2"; there "e"/"E" are hex mantissa digits, so *)
(*) the binary exponent must use "p"/"P" (see spec §6.3). It is the *)
(*) only context in which "p"/"P" is reachable. *)

number      = ["-"] , (int-led | dot-led) , [dec-exponent] ;

(*) Integer-led: one or more digits with an optional fractional part. *)
(*) A trailing dot ("123.") without fractional digits is valid. *)
int-led      = DIGIT , {DIGIT} , [ "." , {DIGIT} ] ;

(*) Dot-led: at least one digit after the dot is required. *)
(*) ".5" and "-.5" are valid; a lone "." is a hard error. *)
dot-led      = "." , DIGIT , {DIGIT} ;

dec-exponent = ("e" | "E") , ["+" | "-"] , DIGIT , {DIGIT} ;

(*) Special IEEE-754 values are bare reserved keywords: nan, inf, and *)
(*) ninf (negative infinity). Like null / true / false / on / off they *)
(*) are lexed as ordinary symbols and reclassified to numeric special *)
(*) values by the validator; a bare word that is not one of these exact *)

```



```

(* spellings stays an ordinary symbol (e.g. "infinity", "nans"). *)
(* There is no sigil and no inline-unit suffix – a unit is supplied *)
(* via the type annotation, e.g. <float:64,m/s> inf. *)
special-number = "nan" | "inf" | "ninf" ;

(* ----- *)
(* 6. Strings *)
(* ----- *)

(* Two or more adjacent quoted literals with only whitespace/comments *)
(* between them are concatenated into a single string token by the *)
(* lexer (string_outro re-enters string_intro on ''). The combined *)
(* byte length is bounded by max_string_length (default UINT16_MAX). *)

string          = string-literal , {ws , string-literal} ;

string-literal  = '' , {string-char} , '' ;

string-char     = safe-string-byte
                  | escape-seq ;

(* safe-string-byte: whitespace 0x09–0x0D, 0x20–0x7E except '' (0x22)*)
(* and '\ ' (0x5C), plus high bytes 0x80–0xFF. *)
(* Control bytes 0x00–0x08, 0x0E–0x1F, and DEL (0x7F) are hard *)
(* errors even inside strings. *)
safe-string-byte = (* byte in [0x09,0x0D] | [0x20,0x21] | [0x23,0x5B] *)
                  (*      | [0x5D,0x7E] | [0x80,0xFF] *) ;

(* Recognised escape sequences (from bvn_escape_lut): *)
escape-seq      = "\" , ( "t"      (* \t = 0x09 HT *)
                          | "n"      (* \n = 0x0A LF *)
                          | "v"      (* \v = 0x0B VT *)
                          | "f"      (* \f = 0x0C FF *)
                          | "r"      (* \r = 0x0D CR *)
                          | "'"      (* \' *)
                          | "\"      (* \\ *)
                          | byte-escape      (* spec 1.1, see below *)
                          | unicode-escape ) ; (* spec 1.1, see below *)

(* All other bytes after "\" are error_illegal_escape_sequence. *)
(* spec 1.1 only – available when the document declares "#!bovnar 1.1" *)
(* or newer; in a 1.0/unversioned document "x"/"u" after "\" stay *)
(* error_illegal_escape_sequence. *)
byte-escape      = "x" , hex-digit , hex-digit ; (* one byte 0xHH *)
unicode-escape   = "u" , "{", hex-digit ,
                  [hex-digit] , [hex-digit] ,
                  [hex-digit] , [hex-digit] , "}" ;

(* 1 to 6 hex digits naming a Unicode scalar. A 7th hex digit before "}" *)
(* (or an empty "{}") is error_illegal_escape_sequence – a structural *)
(* check made before the value is examined. A value that is a surrogate *)
(* (D800..DFFF) or above 10FFFF is error_invalid_codepoint; otherwise the *)
(* scalar is UTF-8 encoded. A \x byte must leave the whole string valid *)
(* UTF-8 (else error_invalid_utf8_byte). *)
hex-digit        = DIGIT | "a".."f" | "A".."F" ;

(* ----- *)
(* 6.5 Inline unit suffix *)
(* ----- *)

(* An inline-unit may follow a scalar number or string value (but not *)
(* a special-number keyword – give its unit via the type annotation). *)
(* At least one whitespace character (ws-mandatory) is the sole *)
(* required separator between the value literal and the unit. *)
(* Inline units attach only at value position: the value rule has a *)
(* scalar-with-unit alternative, while an array element (array-elem) *)
(* uses raw-value, which has no unit option. At the action level the *)

```

```
(* lexer additionally rejects ACT_inline_unit_intro when the *)
(* in_array_element flag is set (a redundant guard). *)
(*)
(*)
(*) INVALID (no whitespace before the unit): *)
(*) .bad = 9.81m/s; (* no separator at all – hard error *) *)
(*)
(*)
(*) Examples: *)
(*) .a = 9.81 m/s; (* space separator *) *)
(*) .b = <float:64> 1.5 k~m; (* with type annotation *) *)
(*) .c = <float:64,s> inf; (* special-number unit via annotation *) *)
(*) .d = 1.5e3 k~J; (* number with exponent *) *)
(*) .e = "FF" m; (* string literal with unit *) *)
(*)
(*) State-machine flow: *)
(*) 1. Whitespace from a number-body state fires ACT_to_number_outro *)
(*) → number_outro. Strings enter string_outro_nosp after the *)
(*) closing "; whitespace transitions to string_outro. *)
(*) Special numbers (nan, inf, ninf) are lexed as symbols and *)
(*) take no inline unit (give the unit via the type annotation). *)
(*) 2. From number_outro / string_outro (whitespace already seen): *)
(*) alpha / "_" / "$" / "%" / "(" / UTF-8 leader (0xC2–0xF4) *)
(*) → ACT_inline_unit_intro *)
(*) 3. ACT_inline_unit_intro pushes the first byte and enters *)
(*) inline_unit_body directly. *)
(*) 4. inline_unit_body accumulates unit characters (including ":" *)
(*) as a valid mid-unit byte) until whitespace (→ *)
(*) inline_unit_outro) or ";" (→ value_outro via finalize). *)
(*)
(*) Inside an array element (in_array_element flag set), *)
(*) ACT_inline_unit_intro immediately rejects with *)
(*) error_unexpected_input_byte. *)
(*)
(*) Semantic rule (enforced by validator): *)
(*) If a type-annotation unit was also given, the inline unit must *)
(*) parse to the identical value_unit_t; mismatch → error_unit_mismatch. *)

inline-unit = inline-unit-start , { inline-unit-char } ;

(* A currency component carries a mandatory "$" sigil (spec 1.0): "$USD", *)
(* "$BTC", or prefixed "<prefix>~$EUR". Bare codes are no longer currencies, *)
(* so a currency code can never collide with a physical-unit symbol. *)
(* "(" lets an inline unit begin with a parenthesised group, e.g. "(m/s)/s", *)
(* so the inline form accepts the same unit grammar as a type-annotation unit. *)
inline-unit-start = ALPHA | "_" | "$" | "%" | "(" | unit-lead-byte ;

inline-unit-char = ALPHA | DIGIT | "_" | "$" | "%" | "+" | "-" | "." | "/" | ":" | "^" |
"~" | "~"
| "(" | ")" | unit-lead-byte | utf8-continuation ;

(* unit-lead-byte: the UTF-8 lead-byte range accepted inside a unit – the *)
(* full 0xC2–0xF4, BROADER than utf8-leader (0xC3–0xF4). Unlike identifier, *)
(* symbol, and reference token starts (where 0xC2 is rejected, see §12), a *)
(* unit may begin with or contain a 0xC2-led code point: the μ prefix *)
(* (U+00B5 = 0xC2 0xB5), the ° / °C units (U+00B0 = 0xC2 0xB0), and the *)
(* "·" product separator (U+00B7 = 0xC2 0xB7). The lexer maps the whole *)
(* 0xC2–0xF4 range to ACT_inline_unit_intro / ACT_copy_inline_unit_byte, *)
(* matching copy_type_byte in a type-annotation unit body. *)
unit-lead-byte = (* byte in [0xC2, 0xF4] *) ;

(*)
(*)
(*) 7. Symbols (unquoted bare-word values in value position) *)
(*)
(*)
(*) A symbol starts with id-start and continues with sym-body-char. *)
(*) Differences from the identifier body action table: *)
(*) "=" is a hard error (not a state transition) *)
```

```

(*) "," / "]" / ";" terminate the symbol cleanly (not errors) *)
(*) "." is a hard error (same as in identifier body) *)

(*) Reserved keywords: a symbol token whose text is exactly "null", *)
(*) "true", "false", "on", or "off" is reclassified by the validator *)
(*) (see null-value / bool-value, §3) and is NOT delivered as a symbol. *)
(*) Any other bare word – including one that merely begins with a *)
(*) keyword, like "ontology" or "nullish" – remains a symbol. *)
symbol      = id-start , {sym-body-char} ;

(*) sym-body-char spans the same byte range as id-body-char; the *)
(*) difference is purely in the termination actions fired for *)
(*) ",", (new_array_value), "]" (array_outro), and ";" (value_outro). *)
sym-body-char = id-body-char ;

(*) _____ *)
(*) 8. References *)
(*) _____ *)

(*) The "&" sigil introduces a dotted path to another key. *)
(*) The stored string starts with "." and includes all dots: *)
(*) "&.foo.bar" stores ".foo.bar". *)
(*) At least one segment is required; "&" alone or "&." without a *)
(*) following id-start character is an error. *)
(*) A reference is stored UNRESOLVED – the path is a string token; the *)
(*) library never dereferences it, so the target need not exist and *)
(*) cycles are not detected. Resolution is the application's job. *)

reference    = "&" , ref-segment , {ref-segment | ref-index} ;
ref-segment  = "." , id-start , {ref-body-char} ;

(*) ref-index (spec 1.1): array indexing in a reference path, e.g. *)
(*) &.matrix[0][1]. The bytes are captured verbatim into the path *)
(*) string (references are stored unresolved); bvn_dom_lookup *)
(*) interprets them at the DOM layer. Gated on a "#!bovnar 1.1" *)
(*) declaration – in a 1.0/unversioned document "[" in a reference is *)
(*) error_unexpected_input_byte. The reference_index sub-state *)
(*) distinguishes a "]" that closes an index from one that closes an *)
(*) enclosing array (which ends the reference). *)
ref-index    = "[" , DIGIT , {DIGIT} , "]" ;

(*) ref-body-char: same byte set as sym-body-char, except that *)
(*) "." (0x2E) is ACT_copy_reference_dot, which pushes the dot into *)
(*) the token buffer and advances to reference_segment_intro for the *)
(*) next segment's id-start character. *)
ref-body-char = sym-body-char | "." ;

(*) _____ *)
(*) 9. Arrays *)
(*) _____ *)

(*) An array is one or more bracket-enclosed rows separated by "/": *)
(*) .a = [1, 2, 3]/[4, 5, 6]; *)
(*) "/" emits ev_array_dim_start; the next "[" opens the next row. *)
(*) Whitespace/comments are permitted on both sides of "/" (between the *)
(*) closing "]" and "/", and between "/" and the next "["). *)
(*) *)
(*) The former adjacent-row syntax "[1,2][3,4]" and the comma-between- *)
(*) rows form "[1,2],[3,4]" are hard errors at the assignment level. *)
(*) array_outro no longer accepts "[" directly; the only continuation *)
(*) for a further row is "/". *)
(*) *)
(*) Nested arrays are valid element values inside a row: *)
(*) [[1,2],[3,4]] is a one-row outer array whose two elements are *)
(*) the inner arrays [1,2] and [3,4]. *)

```

```

(*)
(*) Leading and trailing commas produce null elements:
(*) [1,2,] → four elements: null, 1, 2, null.
(*)
(*) Semantic constraints – two distinct tiers:
(*)
(*) a) Enforced by the LEXER/VALIDATOR (the streaming reader): the
(*) "/"-separated dimension rows of a SINGLE array must all have
(*) the same element count; a mismatch is
(*) error_array_row_size_mismatch (code 2). The lexer checks each
(*) row as it closes, so the error fires at the earliest byte (the
(*) overshooting element, or the closing "]" of a short row).
(*) The per-bracket-pair context save/restore mechanism keeps this
(*) check scoped to one array instance: a "/"-row width never
(*) leaks across a "," boundary or into a sibling branch.
(*) [1,2,3]/[4,5] → error (one array's /-rows: 3 vs 2)
(*) [[1,2]/[3,4,5]] → error (inner /-array's rows: 2 vs 3)
(*)
(*) b) Enforced by the DOM BUILDER ONLY (bvn_dom_parse), NOT by the
(*) streaming reader – `loads`/the SAX callbacks accept these and
(*) deliver every element; only a DOM parse rejects them. Array
(*) elements must be HOMOGENEOUS (spec 1.0): the same kind
(*) throughout, `,`-separated sibling sub-arrays must be
(*) rectangular (match in length), and sibling structs must share
(*) the same shape. A ragged sibling set is
(*) error_array_row_size_mismatch (code 2); a mix of element kinds
(*) (e.g. a scalar beside a sub-array) is
(*) error_array_element_type_mismatch (code 39); sibling structs
(*) with differing shapes is error_struct_shape_mismatch (code 40).
(*) [[1,2],[3,4,5]] → DOM error (ragged sibling sub-arrays)
(*) [1,[2,3]] → DOM error (scalar mixed with a sub-array)
(*) [{.x=1;},{.y=2;}] → DOM error (sibling structs, differing
(*) shapes)
(*) [[1,2]/[3,4], [5,6]/[7,8]] → accepted (same-shape blocks)
(*)
(*) c) Array nesting depth ≤ max_array_nesting (default 64, cap 255).*)

array = array-row , {ws , "/" , ws , array-row} ;

array-row = "[" , ws , [row-content] , ws , "]" ;

(*) row-content is OPTIONAL: "[]" (nothing between the brackets) is an
(*) EMPTY row – zero elements, no ev_data. It is distinct from "[null]"
(*) (one null element) and "[,]" (two null elements). When present,
(*) row-content has at least one element position; leading / trailing /
(*) consecutive commas each expand to a null-value element. An empty row
(*) has width 0, a valid row width for /-dimension consistency.
(*) row-content = array-elem , {ws , "," , ws , array-elem} ;

(*) array-elem has the same structure as a top-level value, except that an
(*) inline unit suffix is not permitted: it uses raw-value (no unit option),
(*) not the value rule's scalar-with-unit alternative (see §6.5).
(*) array-elem = [type-annotation , ws] , raw-value ;

(*) _____
(*) 10. Structs (nested key-value scopes)
(*) _____

(*) A struct "{" ... "}" contains zero or more assignments. "{" is an
(*) empty struct (zero members). After "}" the struct behaves like any
(*) completed value: followed by ";" at the top level, or "," / "]" in
(*) array context.
(*)
(*) Key uniqueness is a DOM-BUILDER constraint, NOT a grammar one: the
(*) lexer/validator (the streaming reader) accepts a repeated key and
(*) delivers BOTH assignments. A DOM parse (bvn_dom_parse) is stricter – *)

```

```

(* a duplicate key within one scope (a struct, or the top-level *)
(* document) is error_duplicate_struct_key. The same key in different *)
(* scopes is always unrelated. *)
(* Nesting depth: struct_nesting_level ≤ max_struct_nesting *)
(* (default 64 when the flag is 0; hard cap 255). *)
(* An unmatched "}" is error_illegal_struct_close. *)

struct      = "{" , ws , {assignment , ws} , "}" ;

(* ----- *)
(* 11. Octet streams (binary-framed escape from the text layer) *)
(* ----- *)

(* A NUL byte (0x00) appearing where a value is expected switches from*)
(* text mode to binary chunk mode. The UTF-8 stream validator is *)
(* suspended for the duration. *)
(* ----- *)
(* Binary protocol (bvn_read_octet_stream): *)
(* Tag 0x01 → chunk follows (2-byte LE length + data bytes) *)
(* Tag 0x00 → end of stream *)
(* Any other tag → error_octet_stream_out_of_sync *)
(* ----- *)
(* After ev_octet_stream_end the lexer enters octet_stream_outro, *)
(* which accepts the same terminators as any other completed value. *)

octet-stream = "\x00" , {os-chunk} , os-end ;

os-chunk      = "\x01" , os-length , os-data ;
os-end        = "\x00" ;

(* 16-bit little-endian length; 0x0000 encodes 65536 bytes. *)
os-length     = byte , byte ; (* little-endian uint16; 0x0000 → 65536 *)

os-data       = {byte} ; (* exactly os-length bytes; not CF *)

(* ----- *)
(* 12. Lexical primitives *)
(* ----- *)

DIGIT         = "0" | "1" | "2" | "3" | "4"
              | "5" | "6" | "7" | "8" | "9" ;

ALPHA         = "A" | "B" | "C" | "D" | "E" | "F" | "G" | "H" | "I" | "J" | "K" | "L" | "M"
              | "N" | "O" | "P" | "Q" | "R" | "S" | "T" | "U" | "V" | "W" | "X" | "Y" | "Z"
              | "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h" | "i" | "j" | "k" | "l" | "m"
              | "n" | "o" | "p" | "q" | "r" | "s" | "t" | "u" | "v" | "w" | "x" | "y" | "z" ;

(* utf8-leader here denotes the range used at identifier, symbol, and *)
(* reference token-start and -body positions: 0xC3–0xF4. *)
(* 0xC2 is the leader for U+0080–U+00BF. It is explicitly overridden *)
(* to ACT_NONE at every identifier, symbol, and reference boundary, *)
(* so code points U+0080–U+00BF cannot start or appear in those *)
(* tokens. 0xC2 IS accepted inside quoted strings and type parameter *)
(* bodies (copy_type_byte), where the full range 0xC2–0xF4 applies. *)
utf8-leader   = (* byte in [0xC3, 0xDF] | [0xE0, 0xEF] | [0xF0, 0xF4] *) ;

utf8-continuation = (* byte in [0x80, 0xBF] *) ;

(* Printable ASCII: 0x20–0x7E *)
printable-byte = (* byte in [0x20, 0x7E] *) ;

(* High byte: any byte 0x80–0xFF. 0xC0, 0xC1, 0xF5–0xFF are *)
(* structurally in this range but rejected by the parallel UTF-8 *)
(* validator (error_invalid_utf8_byte). *)
high-byte     = utf8-continuation | utf8-leader

```

```

        | (* 0xC0 | 0xC1 | 0xF5–0xFF – rejected by UTF-8 check *) ;
byte      = (* any single byte 0x00–0xFF *) ;

```

```

(* _____ *)
(* 13. Constraints not expressible in context-free EBNF *)
(* _____ *)

```

```

(*)

```

a) UTF-8 validity

Every byte in the text layer (outside octet-stream regions) must form valid UTF-8 when fed through `bvn_utf8_feed`. The UTF-8 validator runs in parallel with the action state machine; a non-UTF-8 byte causes `error_invalid_utf8_byte` regardless of which lexer state the parser is in. Overlong encodings and surrogate halves are rejected.

b) BOM uniqueness

The UTF-8 BOM (EF BB BF) is legal only at byte offset 0. If the stream begins with "#" instead (a first-line comment), the special `first_comment *` state machine watches the comment body for the exact byte sequence 0xEF 0xBB 0xBF and rejects it as `error_invalid_byte_order_mark`. This detection applies only while the `first_comment *` states are active (i.e. until the first newline); subsequent regular comments (`comment_intro` state) have no BOM detection, so U+FEFF encoded as 0xEF 0xBB 0xBF appearing in any later comment body or string literal is valid UTF-8 and is accepted.

c) 0xC2 at token boundaries

The byte 0xC2 (UTF-8 leader for U+0080–U+00BF) is mapped to `ACT_NONE` at identifier, symbol, and reference token-start and token-body positions, forbidding code points U+0080–U+00BF in those tokens. 0xC2 is accepted inside quoted strings and in type parameter bodies (`copy_type_byte` / `type_body_outro`).

d) Special-number keywords

The special floats `nan` / `inf` / `ninf` are bare reserved keywords (no sigil). They are lexed as symbols and reclassified to numeric special values by the validator, like `null` / `true` / `false` / `on` / `off`; any other bare word (e.g. "infinity", "nans") stays a symbol. They take no inline unit – a unit is supplied via the annotation.

e) String concatenation

Adjacent string literals with only whitespace/comments between them are concatenated into a single token by the lexer. The combined byte length is bounded by `max_string_length` (default `UINT16_MAX` = 65535 bytes).

f) Number exponent in bare literals vs. quoted strings

Bare number literals only accept "e"/"E" as exponent markers (`exp_intro` state).

g) Array row-size consistency (LEXER/VALIDATOR – the streaming reader)

The "/"-separated dimension rows of a single array must all have the same element count; the lexer checks each row as it closes, so the violation is reported at the earliest offending byte. The per-bracket-pair context save/restore mechanism keeps the /-row check scoped to one array instance.

Violation: `error_array_row_size_mismatch`

(e.g. `[1,2,3]/[4,5]`, or `[[1,2]/[3,4,5]]`).

Sibling homogeneity (DOM BUILDER ONLY – `bvn_dom_parse`; the streaming reader accepts these). Array elements must be the same kind throughout, `,`-separated sibling sub-arrays must be rectangular (match in length, so a bare array/matrix is rectangular), and sibling structs must share

the same shape. A ragged sibling set is `error_array_row_size_mismatch` (e.g. `[[1,2],[3,4,5]]`); a mix of element kinds is `error_array_element_type_mismatch` (code 39, e.g. `[1,[2,3]]`); sibling structs with differing shapes is `error_struct_shape_mismatch` (code 40, e.g. `{x=1;},{y=2;}`). `loads()`/the SAX callbacks deliver such arrays unchanged; only a DOM parse rejects them.

- h) Array and struct nesting limits
 - Array nesting (bracket depth): `max_array_nesting` (default 64 when the flag is 0; hard cap 255).
 - Struct nesting: `max_struct_nesting` (default 64 when the flag is 0; hard cap 255).
- i) Type-annotation / value compatibility (enforced by validator)
 - "utf8" requires a string value.
 - "uint", "sint", "float", "float_fix", "float_dec" accept a number or string value; a quoted string allows special float notation.
 - "utf8" and "bool" are parameterless: the lexer accepts the ":" and parameter bytes, but `bvn_parse_type_annotation` rejects any width, base, q, or unit parameter as `error_illegal_value_type` (e.g. "`<utf8:8>`" and "`<utf8:m>`" are both errors).
 - "uint" rejects negative number literals (`error_value_out_of_range`).
 - Digit characters in the value must be valid for the declared base (`error_digit_not_in_base`).
 - The numeric value must fit in the declared bit-width (`error_value_out_of_range`).
 - For "float": accepted widths are 0, 16, or any multiple of 32 up to `BVN_FLOAT_MAX_PREC` (32768); accepted bases are 10 (or absent) and 16; any other base → `error_illegal_value_type`.
 - For "float_fix": accepted widths are 0, 16, 32, 64, 128, 256; the base parameter (`_N`) is forbidden (`error_illegal_value_type`); the Q parameter (`qN`) must satisfy $0 \leq Q < \text{effective_width}$ (`error_illegal_value_type` if violated); `Q` stored in `value_type_spec_t.base`; retrieved via `bvn_effective_q`.
 - For "float_dec": accepted widths are 0, 16, 32, 64, 128, 256; the base parameter (`_N`) is forbidden (`error_illegal_value_type`); the Q parameter is not valid for `float_dec`.
 - Neither "float_fix" nor "float_dec" can be synthesised by the validator's default type synthesis; both require explicit annotation.
- j) Identifier (key) non-empty
 - The grammar requires at least one id-start character after ".".
 - An empty key ("=") triggers `error_empty_identifier`.
- k) Struct closing brace without matching open
 - A "}" seen when `struct_nesting_level == 0` is `error_illegal_struct_close`.
- l) Comma in non-array context
 - "," triggers `ACT_new_array_value`, which checks `array_nesting_level > 0` and returns `error_unexpected_input_byte` if the condition is not met. A stray "," at the top level is therefore an error.
- m) Octet stream out-of-sync
 - Any byte other than `0x01` (chunk) or `0x00` (end) as the tag byte inside an octet stream causes `error_octet_stream_out_of_sync`.
- n) Error recovery (continue_on_error mode)
 - When `bvnr_read_flags_t.continue_on_error` is set, any parse error invokes the `on_error` callback and then enters the resync state machine (states: `resync`, `resync_string`, `resync_string_escape`, `resync_comment`). The resync machine skips bytes while tracking bracket and brace nesting depths, and attempts to re-synchronise at the next ";" at the current nesting depth. `recovery_count` (`bvnr_reader_get_recovery_count`) is incremented immediately when an error triggers entry into resync mode, not when resync

completes at ";". EOF inside any resync state propagates the original error code rather than overwriting it with `error_got_incomplete_bvnr_stream`.

o) Inline unit suffix

An inline-unit follows a scalar number or string in non-array context only. A special-number keyword (`nan`, `inf`, `ninf`) takes NO inline unit – it is lexed as a symbol, and symbols accept no inline suffix; `"inf m/s"` is `error_unexpected_input_byte`. Supply a unit for a special value through the type annotation (`<float:64,m/s> inf`). At least one whitespace character MUST separate the value literal from the unit. The only accepted form is `"value unit"`; any other form (no separator or colon) is a hard error.

State-machine flow:

1. Whitespace from number-body states (`zero_intro`, `copy_number_byte`, `fraction_intro`, `copy_fraction_byte`, `copy_exp_byte`) fires `ACT_to_number_outro` → `number_outro`. Strings enter `string_outro_nosp` after the closing quote; whitespace transitions to `string_outro`. A comment's terminating newline also counts as whitespace (`comment_outro` promotes `string_outro_nosp` → `string_outro`).
2. From `number_outro` / `string_outro` (at least one ws consumed):
`alpha` / `"_"` / `"$"` / `"%"` / `"("` / UTF-8-leader (`0xC2-0xF4`)
 → `ACT_inline_unit_intro`
3. `ACT_inline_unit_intro` pushes the first byte and enters `inline_unit_body` directly.
4. `inline_unit_body` accumulates unit characters (including `":"` as a valid mid-unit byte) until whitespace (→ `inline_unit_outro`) or `;"` (→ `value_outro` via `finalize`).

Inside an array element (`in_array_element` flag set), `ACT_inline_unit_intro` immediately rejects with `error_unexpected_input_byte`.

The validator's `bvn_val_receive` calls `bvn_parse_unit` on the inline buffer and compares the result to `v->parsed_unit` using `memcmp` on the full `value_unit_t` structure (valid because both structs are zero-initialised before parsing and all padding is zeroed). When `has_annotation_unit` is true and the comparison fails, the validator returns `error_unit_mismatch`.

When no annotation unit was present, the inline unit is silently adopted as the effective unit.

*)